Complete SCF calculations on 1 to 16 GPUs: time, GPU memory and host memory

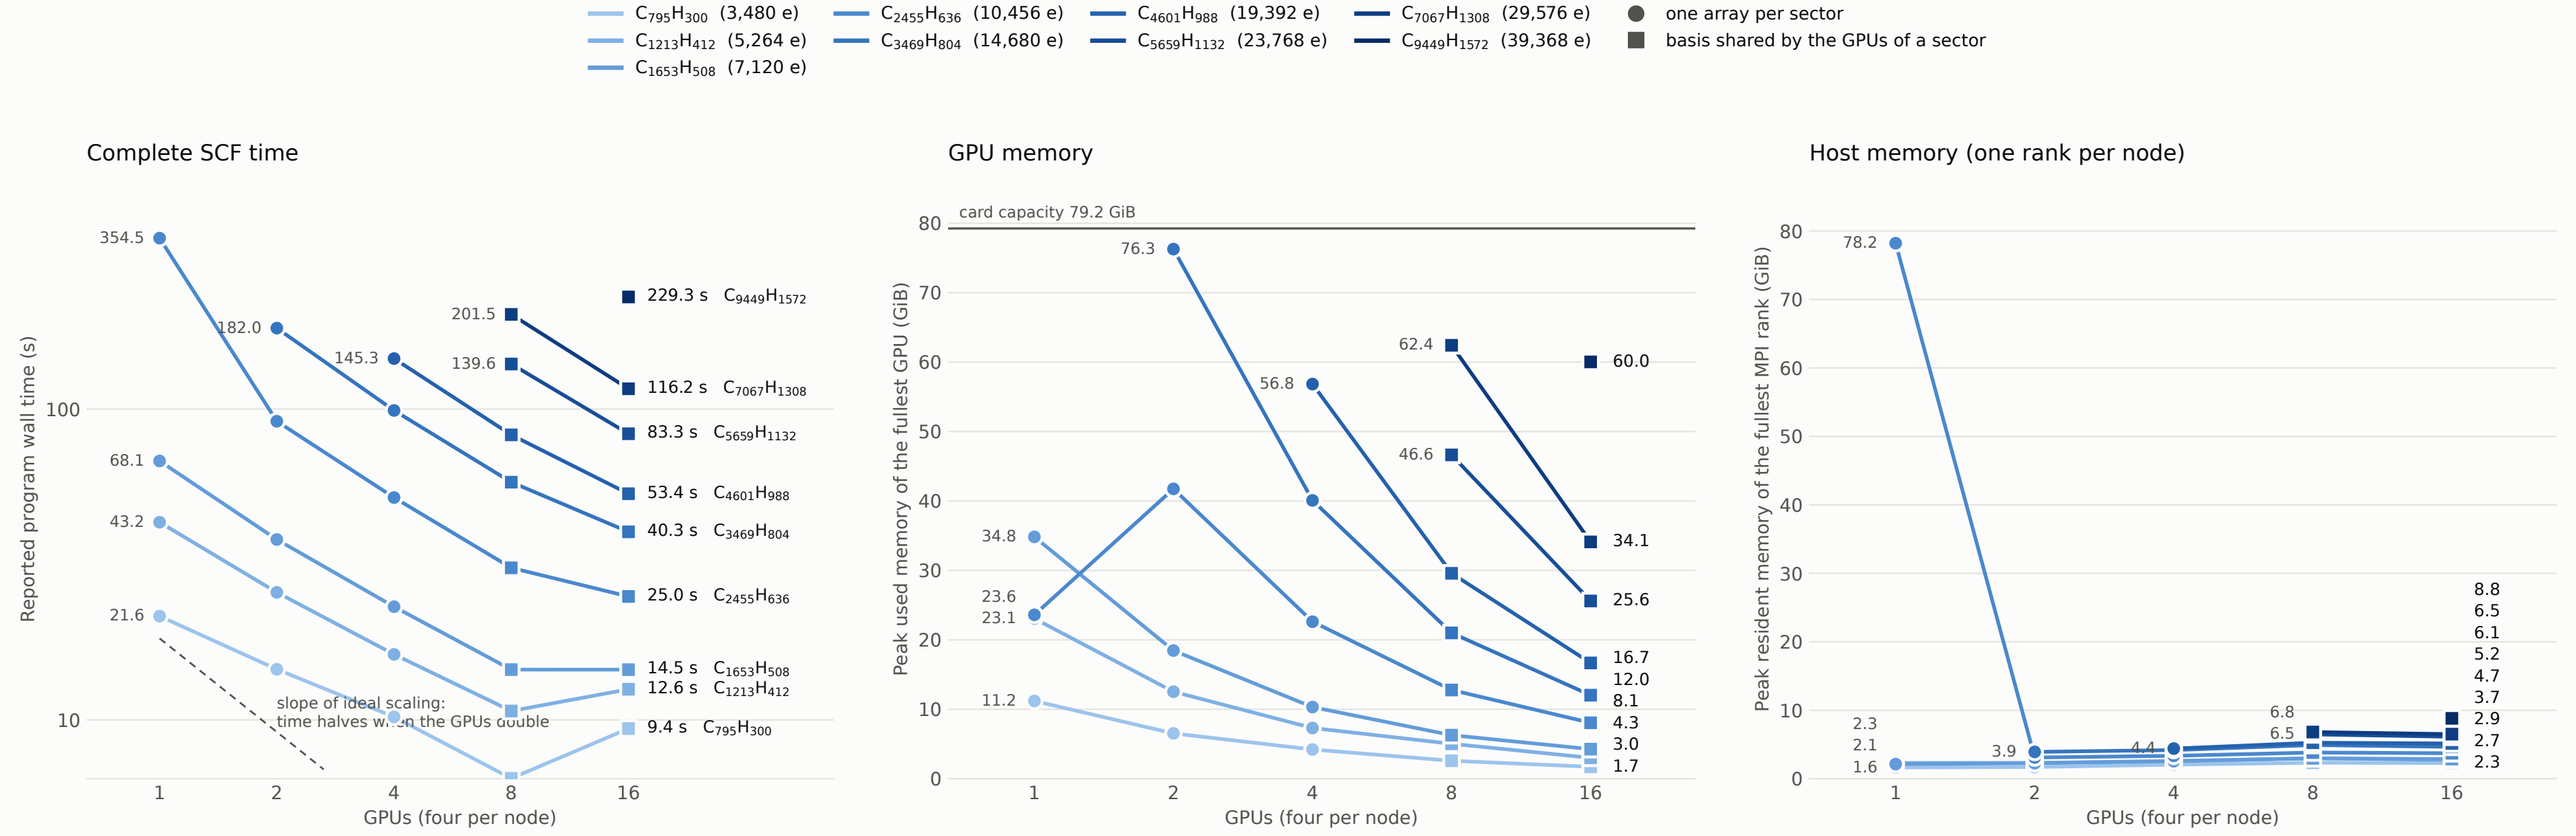

One measured run per point after a warm-up (C2455H636 on 1 GPU, C3469H804 on 2 GPUs: one run); totals identical to series F11 to the last bit, same SCF steps. GPU memory: nvidia-smi samples every 200 ms, fullest GPU, including contexts and allocator caches. Host memory: resident high-water mark of the fullest MPI rank. Series F12 (5 A of vacuum): code defaults, double precision, first calculations without symmetry cache. Source 5fc7363.